Week 8

# Unsupervised learning - introduction : Clustering

## Unsupervised learning

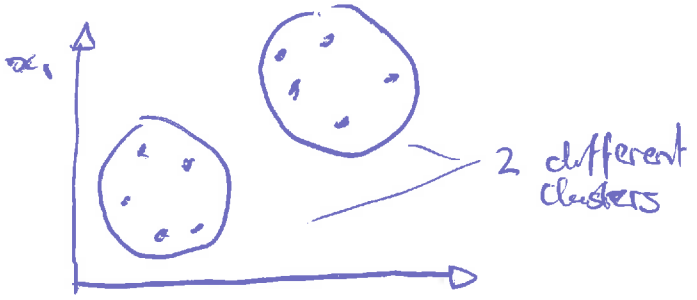


Training set: $\{x^{(1)}, x^{(2)}, x^{(3)} \dots x^{(m)}\}$

## Applications of clustering

- Market segmentation
- Social network Analysis
- Organise computing clusters
- astronomical Analysis

# K-means Algorithm

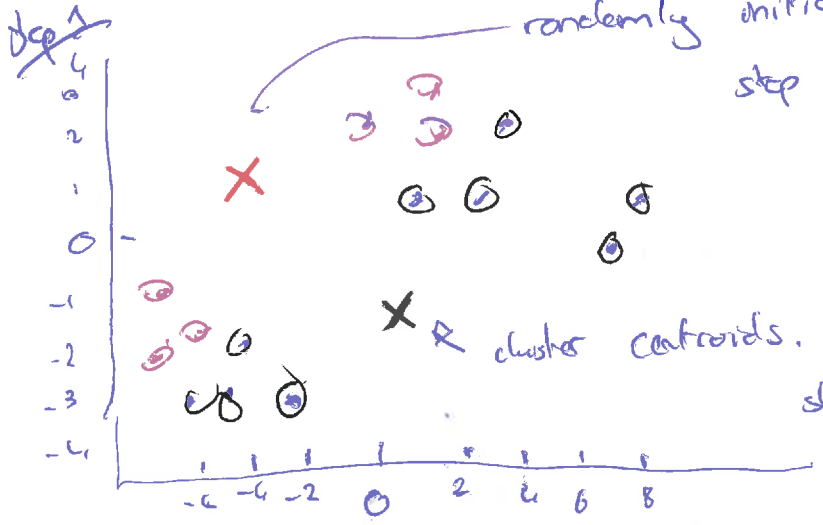


step 1) Cluster assignment step.

So if a data point is closer to a cluster centroid → it is assigned to that cluster centroid.

step 2) move centroid step

→ compute average of all the assigned datapoints of a cluster and move the centroid to this new centrepoint.

→ after awhile the centroid locations no longer change.

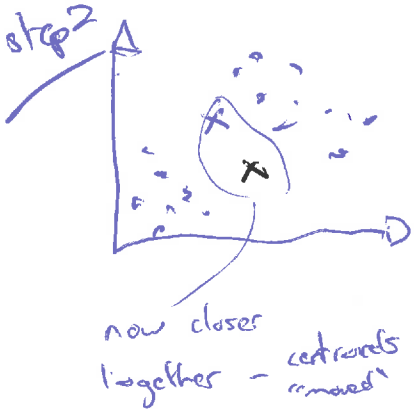

# K-means algorithm

Input:
- $K$ (number of clusters) ✓
- training set $\{x^{(1)}, x^{(2)}, \ldots\}$ ✓ (no $x_0$...)

$x^{(i)} \in \mathbb{R}^n$ (drop $x_0 = 1$ convention)

vasanmamane@gmail.com

## K-means algorithm

Randomly initialize $K$ cluster centroids $\mu_1, \mu_2, \ldots, \mu_K \in \mathbb{R}^n$ (total number of centroids)

Repeat {

cluster assignment step:
for $i = 1$ to $m$
$c^{(i)}$ := index (from 1 to $K$) of cluster centroid closest to $x^{(i)}$
$\min_k \|x^{(i)} - \mu_k\|^2$ ← length.

move centroid step:
for $k = 1$ to $K$
$\mu_k$ := average (mean) of points assigned to cluster $k$
Ex: let assume this cluster ($k=2$) has the values $x^{(1)}, x^{(5)}, x^{(6)}, x^{(10)}$ assigned to it.
$\therefore c^{(1)} = 2, c^{(5)} = 2, c^{(6)} = 2, c^{(10)} = 2.$

$$\mu_2 = \frac{1}{4}\left[x^{(1)} + x^{(5)} + x^{(6)} + x^{(10)}\right) \in \mathbb{R}^n$$

}

- cluster assignment step: distance between $i^{th}$ training example and cluster centroid
So $\therefore$ picking the cluster centroid $c^{(i)}$ that is closest to our training example $x^{(i)}$

notes: delete a cluster if it has no points assigned to it
or reinitialize it through randomization

## K-means for non-separated Clusters

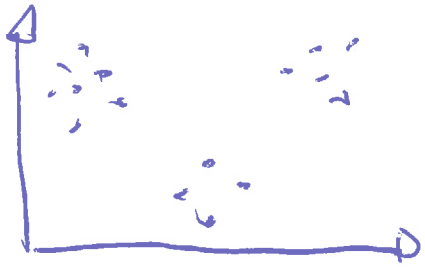

T-shirt sizing S, M, L

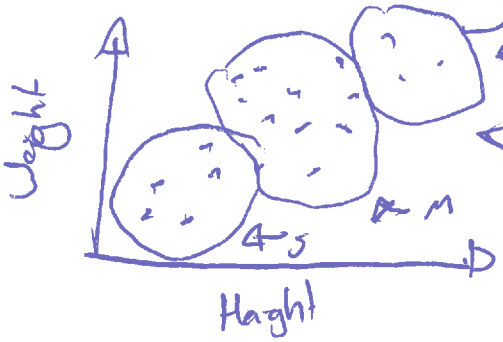


running K means to search for 3 clusters for S, M, L.

Week 8

# Clustering - optimization objective

## K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \ldots, K)$ to which example $x^{(i)}$ is currently assigned.

$\mu_k$ = cluster centroid $k$ ($\mu_k \in \mathbb{R}^n$)

lower case $k$ is an index into $k = \{1, 2, \ldots K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which $x^{(i)}$ has been assigned

$x^{(i)} \rightarrow 5 \therefore c^{(i)} = 5$

$x$ has been assigned to cluster 5

$\mu_{c^{(i)}} = \mu_5 = ?$

Optimization objective:

$$J(c^{(1)}, \ldots, c^{(m)}, \mu_1, \ldots, \mu_K) = \frac{1}{m} \sum_{i=1}^{m} \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$\mu_{c^{(i)}}$ = location of the cluster centroid to which $x^{(i)}$ has been assigned

$$\min_{c^{(1)}, \ldots, c^{(m)}, \mu_1, \ldots, \mu_K} J(c^{(1)}, \ldots, c^{(m)}, \mu_1, \ldots, \mu_K)$$

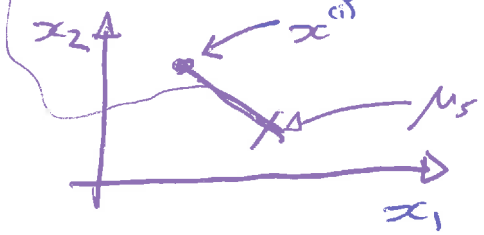


→ Sometimes called the "distortion cost function"

step 1 → The cluster assignment step in minimizing $J(\ldots)$ with respect to $c^{(1)}, c^{(2)}$ etc $\ldots c^{(m)}$ (holding $\mu_1, \ldots, \mu_K$ fixed)

step 2 → minimizes $J(\ldots)$ w.r.t. $\mu_1 \ldots \mu_K$.

## Clustering - Random initialization

should have $K < m$

Randomly pick $K$ training examples

Set $\mu_1, \dots, \mu_K$ equal to these $K$ examples

$\mu_1 = x^{(i)}$

$\mu_2 = x^{(j)}$

$\vdots$

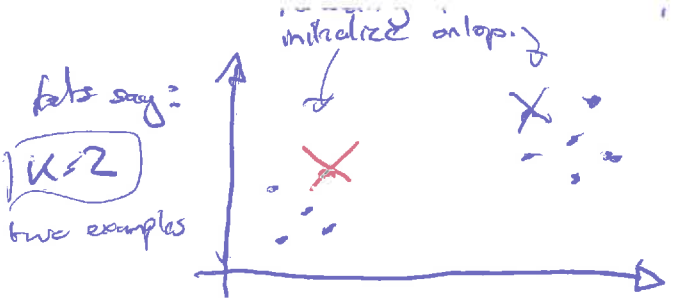


K means can converge to different solutions depending how the clusters were initialized

---

Local optima. → can get stuck at different local optima.

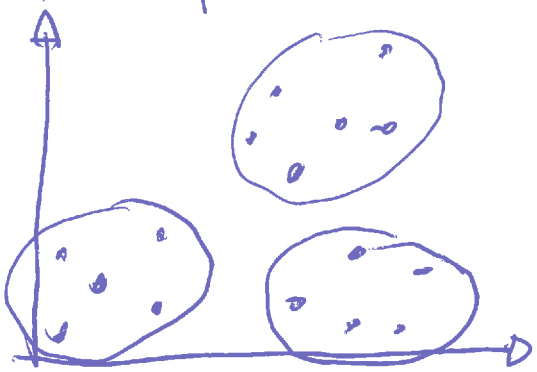

Initialize K means lots of times (randomly) to increase the chances of finding the global minimum.

---

### Random initialization

typical 50 → 1000

for i = 1 to 100 {

Randomly initialize K-means

Run K means to get $c^{(1)}, \dots, c^{(m)}, \mu_1 \dots \mu_K$

Compute cost function (distortion)

$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots \mu_K)$

Pick clustering that gives lowest cost $J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots \mu_K)$

$K = 2 \rightarrow$ ~~100~~ 10

If you have $K > 10$, multiple random initializations will not make much difference

Week 8

# Clustering - choosing the number of clusters

## What is the right value of K?

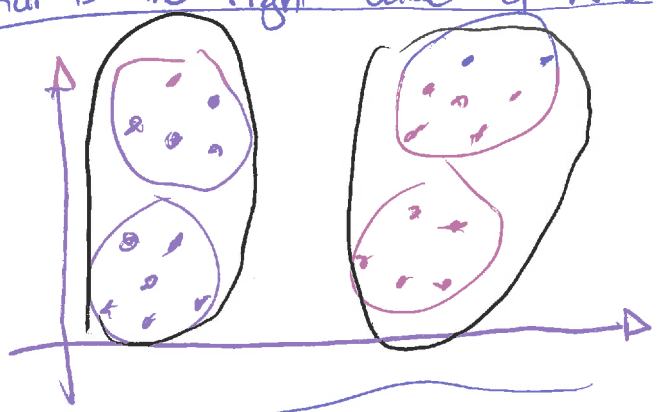

## Choosing the value of K

Elbow method: "worth a shot"

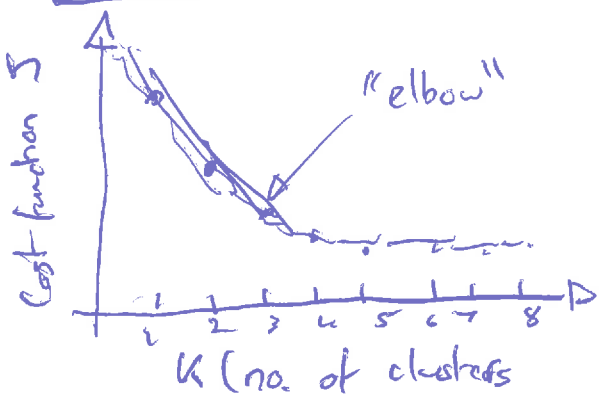


Sometimes this happens

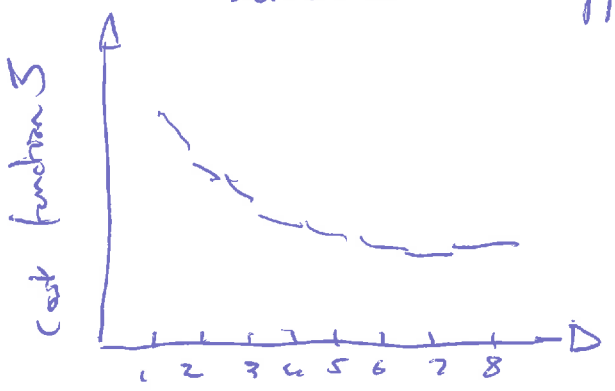


## Choosing the value of K

Sometimes, you're running K means to get clusters for some later/ downstream purpose. Evaluate K-means based on a metric for how well it performs for that later purpose.

K=3 S, M, L T-shirt size

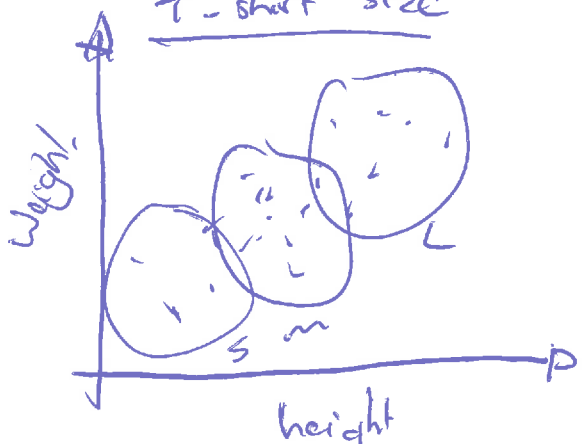


K=5 T-shirt size

XS, S, M, L, XL

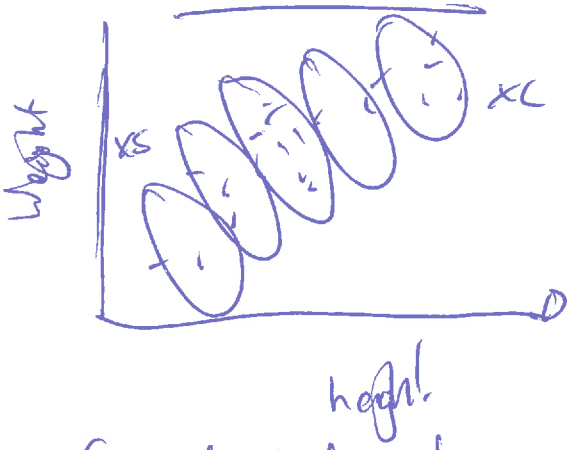


deciding what set of sizes for t shirts to manufacture is

Week 8

# Dimensionality Reduction – Motivation I: data compression

## Data compression

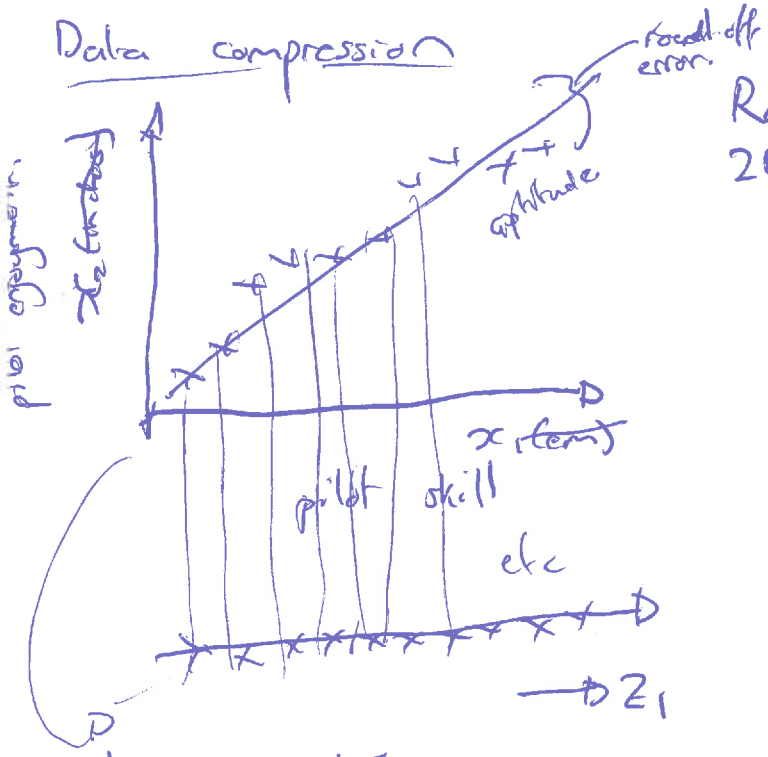


Reduce data from 2D to 1D

$x^{(1)} \in \mathbb{R}^2 \rightarrow z^{(1)} \in \mathbb{R}$

$x^{(2)} \in \mathbb{R}^2 \rightarrow z^{(2)} \in \mathbb{R}$

$\vdots$

$x^{(m)} \rightarrow z^{(m)}$

need only 1 real number to specify a point on a line.

---

## Data Compression

1000D to 100D

reduce data from 3D to 2D.

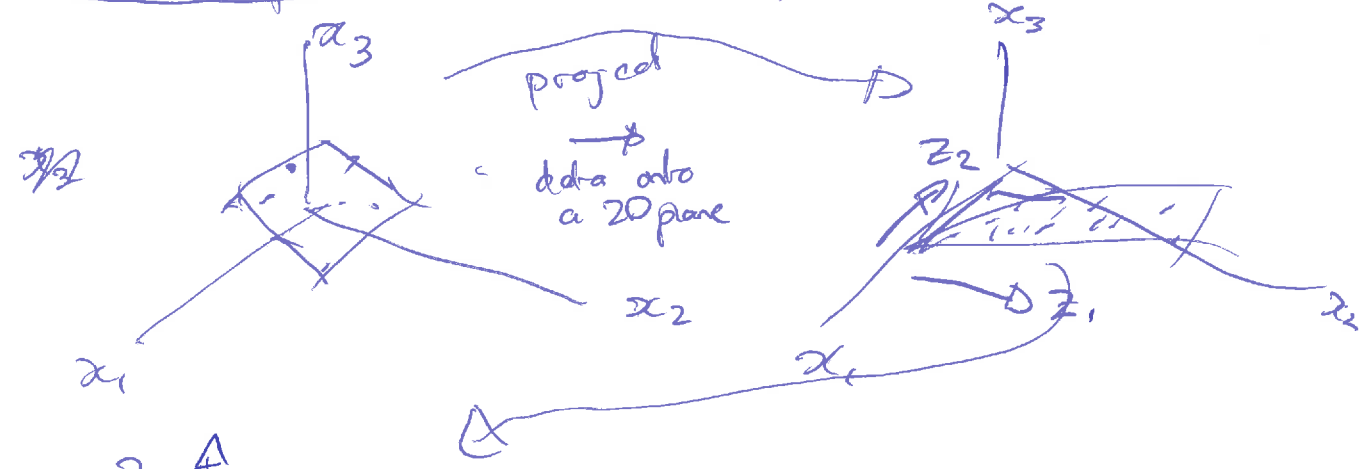


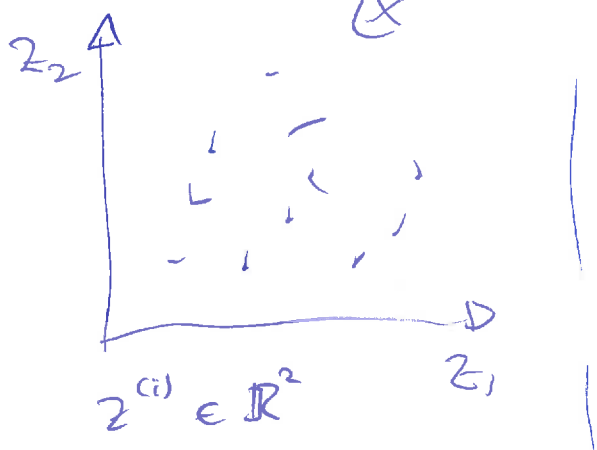


$z^{(i)} \in \mathbb{R}^2$

$z = \begin{bmatrix} z_1 \\ z_2 \end{bmatrix}$ $\quad z^{(i)} = \begin{bmatrix} z_1^{(i)} \\ z_2^{(i)} \end{bmatrix}$

# Dimensionality Reduction 3 Motivation II 2 data visualization

Reduce data from 50D to 2D.

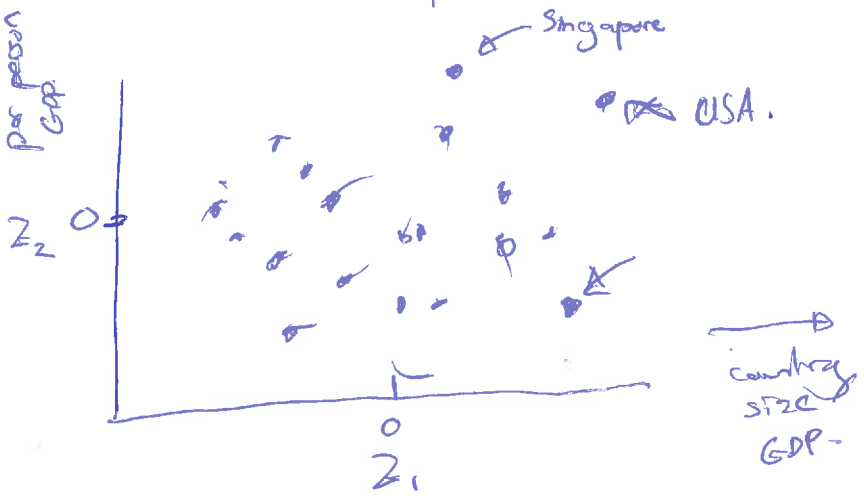

# Dimensionality Reduction - principal component analysis algorithm

## Data preprocessing

Training set: $x^{(1)}, x^{(2)}, \ldots, x^{(m)}$

preprocessing (feature scaling / mean normalization)

$$\mu_j = \frac{1}{m}\sum_{i=1}^{m} x_j^{(i)}$$

Replace each $x_j^{(i)}$ with $x_j - \mu_j$

If different features on different scales (e.g. $x_1$ = size of house, $x_2$ = number of bedrooms), scale features to have comparable range of values.

$$x_j^{(i)} \leftarrow \frac{x_j^{(i)} - \mu_j}{s_j}$$

$s_j$ ← standard deviation of feature $j$.

## Principal Component Analysis Algorithm

- Reduce data from 2D to 1D
- Reduce data from 3D to 2D.

$x^{(i)} \in \mathbb{R}^2 \rightarrow z^{(i)} \in \mathbb{R}$

$z = \begin{bmatrix} z_1 \\ z_2 \end{bmatrix}$

e.g.

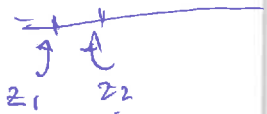


## Principal Component Analysis (PCA) Algorithm

Reduce data from n-dimensions to k-dimensions

Compute "covariance matrix"

not to be confused with summation....

$$\Sigma = \frac{1}{m}\sum_{i=1}^{n} (x^{(i)})(x^{(i)})^T$$

$x^{(i)}$: $n \times 1$; $(x^{(i)})^T$: $1 \times n$ → $n \times n$ matrix

Compute "eigenvectors" of matrix $\Sigma$ → $n \times n$ matrix

→ `[U, S, V] = svd(Sigma);` (or `eig(sigma)`)

↑ octave ↑ singular value decomposition

$$U = \begin{bmatrix} | & | & | & & | \\ u^{(1)} & u^{(2)} & u^{(3)} & \cdots & u^{(n)} \\ | & | & | & & | \end{bmatrix}$$

$U \leftarrow \mathbb{R}^{n \times n}$

$u^{(1)}, \ldots, u^{(k)}$ ← k direction onto which we want to project the data.

K

vectors for plane

# Dimensionality reduction - principal component analysis algorithm

from $[U, S, V] = svd(Sigma)$, we get:

$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \cdots & u^{(n)} \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times n}.$$

$K$.

$x \in \mathbb{R}^n \rightarrow z \in \mathbb{R}^k$

$$z = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \cdots & u^{(k)} \\ | & | & & | \end{bmatrix}^T x^{(i)} = \begin{bmatrix} — & u^{(1)} & — \\ — & u & — \\ — & (u^{(k)})^T & — \end{bmatrix} x^{(i)}$$

$n \times k$. $\quad k \times n \quad n \times 1$

$U_{reduce}$

$k \times 1$

$z \in \mathbb{R}^k$

---

## PCA Algorithm summary

$$Sigma = \frac{1}{m} \sum_{i=1}^{m} (x^{(i)})(x^{(i)})^T$$ implemented in octave looks like → `sigma = (1/m) * X' * X;`

```
[U, S, V] = svd(Sigma);
Ureduce = U(:, 1:k);
z = Ureduce' * x;
```

octave code

$x \in \mathbb{R}^n$ ~~$x_0 = 1$~~

# Dimension Reduction: Reconstruction from compressed representation

In the previous examples we reduced a 2D (or $N^{th}$ D) dimension space, down to 1D using

$(x_2, x_1)$ $\left\{ z = U_{reduce}^T x \right.$ ← new dimension (1D)

If we reverse it, we need to do the following:

$\left\{ z \in \mathbb{R} \rightarrow x \in \mathbb{R}^2. \right.$

$x_{approx} = \underbrace{U_{reduce}}_{n \times k} \cdot \underbrace{z}_{k \times 1}$

$\underbrace{\quad}_{n \times 1}$

$\simeq x$ ← original $x$.

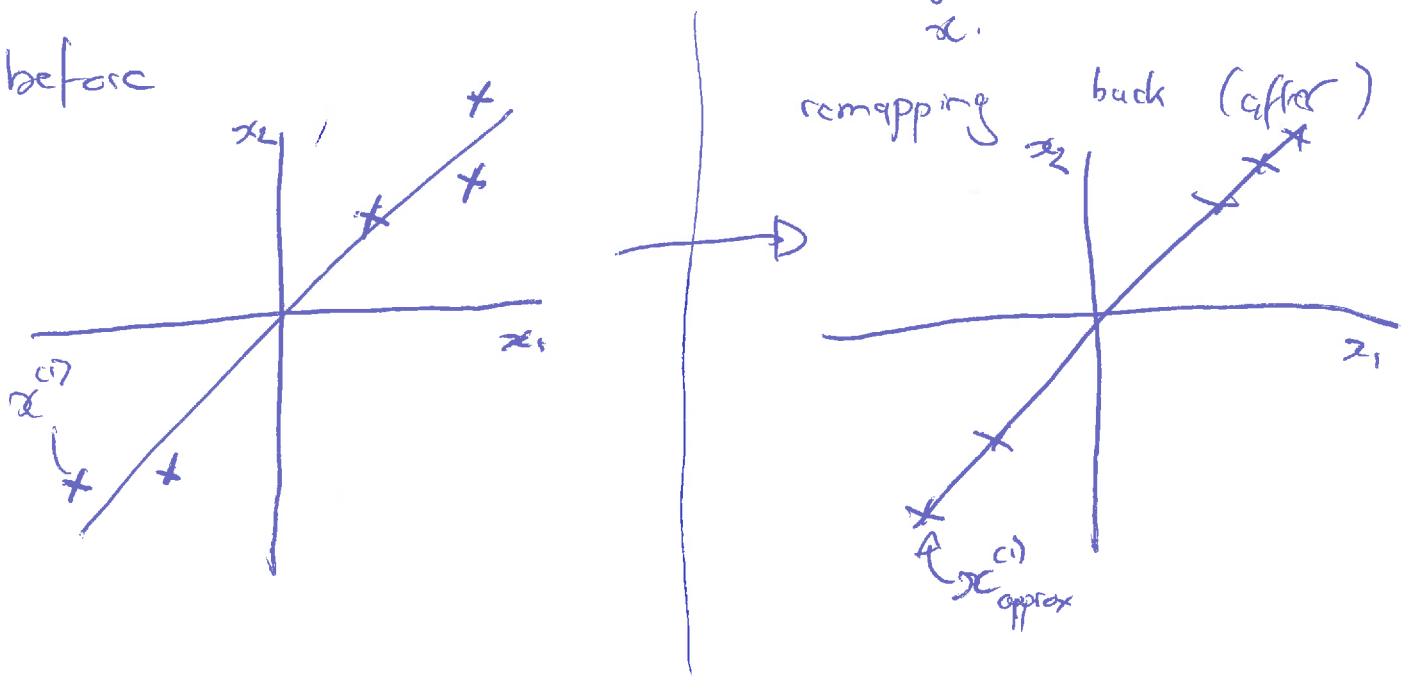

# Dimensionality Reduction – choosing the number of principal components

## Choosing $k$ (number of principal components)

Average squared projection error: $\frac{1}{m}\sum_{i=1}^{m} \| x^{(i)} - x_{approx}^{(i)} \|^2$

Total variation in the data: $\frac{1}{m}\sum_{i=1}^{m} \| x^{(i)} \|^2$

Typically, choose $k$ to be smallest value so that

$$\frac{\frac{1}{m}\sum_{i=1}^{m} \| x^{(i)} - x_{approx}^{(i)} \|^2}{\frac{1}{m}\sum_{i=1}^{m} \| x^{(i)} \|^2} \leq 0.01 \quad (1\%)$$

or $0.05$ (5%).

"99% of variance is retained" / 95% of variance retained

## Choosing $k$ (number of principal components)

Algorithm:

Try PCA with $k = 1$, ~~$k=2$~~ (if not less than 0.01, try 2 if not try 3) $k = 3$ etc

Compute $U_{reduce}, z^{(1)}, z^{(2)}, \ldots, z^{(m)}, x_{approx}^{(1)}, \ldots, x_{approx}^{(m)}$

Check if:

$$\frac{\frac{1}{m}\sum_{i=1}^{m} \| x^{(i)} - x_{approx}^{(i)} \|^2}{\frac{1}{m}\sum_{i=1}^{m} \| x^{(i)} \|^2} \leq 0.01$$

$[U, S, V] = svd(sigma)$

$S = \begin{bmatrix} S_{11} & & & 0 \\ & S_{22} & & \\ & & S_{33} & \\ & & & \ddots \\ 0 & & & S_{nn} \end{bmatrix}$

For given value $k$ ($k \leq 3$)

$$1 - \frac{\sum_{i=1}^{k} S_{ii}}{\sum_{i=1}^{n} S_{ii}} \leq 0.01 \quad (n = n)$$

## Summary

→ Pick smallest value of $k$ for which

$$\frac{\sum_{i=1}^{k} S_{ii}}{\sum_{i=1}^{m} S_{ii}} \geq 0.99 \quad (99\% \text{ variance retained})$$

"Choose what dimension to reduce the data too"

# Dimensionality reduction - advice for applying PCA

## Supervised learning speedup

$(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$

$x^{(i)} \in \mathbb{R}^{10000}$

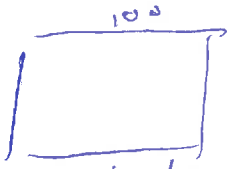


Extract inputs:

Unlabeled dataset: $x^{(1)}, x^{(2)}, \dots, x^{(m)} \in \mathbb{R}^{10000}$

↓ PCA (mapped to $z$)

$z^{(1)}, z^{(2)}, \dots, z^{(m)} \in \mathbb{R}^{1000}$

New representation

New training set: $(z^{(1)}, y^{(1)}), (z^{(2)}, y^{(2)}), \dots, (z^{(m)}, y^{(m)})$

if using logistic regression for example:

$$h_\theta(z) = \frac{1}{1 + e^{-\theta^T z}}$$

Note: Mapping $x^{(i)} \to z^{(i)}$ should be defined by running PCA only on the training set. This mapping can be applied as well to the examples $x^{(i)}_{cv}$ and $x^{(i)}_{test}$ in the cross validation and test sets.

Reduce parameter learned by PCA → fit parameters only to training set. not cv or test set. → apply mapping to train, cv & test sets

## Application of PCA

- Compression
  - Reduce memory/disk needed to store data.
  - Speed up learning algorithm
  - choose $k$ by % of variance retained.
- Visualization ✓
  - $k=2$ or $k=3$ ✓

# Dimensionality reduction – advice for applying PCA

## Bad use of PCA: To prevent overfitting:

Use $z^{(i)}$ instead of $x^{(i)}$ to reduce the number of features $k < n$ (1000 – 1000)

Thus, fewer features, less likely to overfit. Bad!

This might work ok, but it is not a good way to address overfitting. Use regularization instead.

(λ ftw. (λ knows what the values of y are ...) ← minimization term)

(PCA)
→ it throws away data, without knowing what the values of y are.

---

PCA is sometimes used where it should not be.

Design of ML system:

→ Get training set $\{(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})\}$

~~→ Run PCA to reduce $x^{(i)}$ in dimension to get $z^{(i)}$~~

→ Train logistic regression on $((z^{(1)}, y^{(1)}), \dots (z^{(m)}, y^{(m)})$

→ Test on test set: Map $x^{(i)}_{test}$ to $z^{(i)}_{test}$. Run $h_\theta(z)$ on $(z^{(1)}_{test}, y^{(1)}_{test} \dots z^{(m)}_{test}, y^{(m)}_{test})$

→ How about the whole thing without PCA

→ Try running the normal thing with original/raw data $x^{(i)}$. If the implementation is too slow etc try implementing $z^{(i)}$ etc