

Week 11

Application example: photo OCR. \rightarrow problem description and pipeline

Photo OCR pipeline

\rightarrow sliding windows

1. text detection. (1-5 engineers)
2. character segmentation. (1-5 engs)
3. character classification (1-5 etc)
4. (spelling correction)

Sliding Windows

Supervised learning for pedestrian detection

x = pixels in 82×36 image patches

positive examples ($y=1$)

negative examples ($y=0$)

1) run patch through classifier

2) stride rectangle over a bit. (shifting amount (step-size/stride))


3) increase image patch size \rightarrow resize down to 82×36

\rightarrow take output of classifier \rightarrow expand it; if a character is detected within the vicinity of another ~~character~~ character \rightarrow can conclude that it is text
use heuristics to look for "long rectangles"

character segmentation \rightarrow look for splits

splits = $y=1$

no split $y=0$



Week 11

Application example: photo OCR: getting lots of data!
artificial data synthesis

one of the best things to do in ML.

take a low bias algorithm, \rightarrow train on a massive data set.

synthesize data - by making more examples using real facts
- distort existing examples to create more data

can do the same with audio seg.

- distortion introduced should be representation of the type of noise / distortion in the test set.

- Usually does not help to add purely random / meaningless noise to your data

Discussion on getting more data

1) Make sure you have a low bias classifier before expending effort.
(Plot learning curves) e.g. keep increasing the number of features / number of hidden units in neural network until you have a low bias classifier.

2) How much work would it be to get 10x as much data as we currently have?

- artificial data synthesis

- collect / label it yourself.

- "crowd source"

↳ e.g. amazon mechanical turk.


how many hours will it take to collect 10x as many examples
 $n = 1000 \rightarrow 10,000$
10 seconds / example

Week 11

Application example: photo OCR.

: Ceiling analysis, what part of the pipeline to work on next?

estimating the errors due to each component (ceiling analysis)



what part of the pipeline should you spend the most time trying to improve? or where should you allocate scarce resources?

Component	Accuracy	
Overall system	73%	
text detection	89%	17%
character seg.	90%	17%
character recog.	100%	10%

provide it the correct text detection output

Another ceiling analysis example

face recognition from images
(artificial example)

